



# Introduction to Neural Networks



# Deep Learning

- This series of lectures will cover key theory aspects:
  - Neurons and Activation Functions
  - Cost Functions
  - Gradient Descent
  - Backpropagation



# Deep Learning

- Once we build a general high level understanding we will code out all these topics manually with Python, without the use of a deep learning library.
- Then we can move on to using TensorFlow!



# Deep Learning

- Understanding a high level overview of these key elements will make it much easier to understand what is happening when we begin to use TensorFlow!
- Tensorflow has direct connections to these concepts in its syntax!



# Let's get started!



# Introduction to the Perceptron



# Deep Learning

- Before we launch straight into neural networks, we need to understand the individual components first, such as a single “neuron”.



# Deep Learning


- Artificial Neural Networks (ANN) actually have a basis in biology!
- Let's see how we can attempt to mimic biological neurons with an artificial neuron, known as a perceptron!





# Deep Learning


- The biological neuron:





# Deep Learning

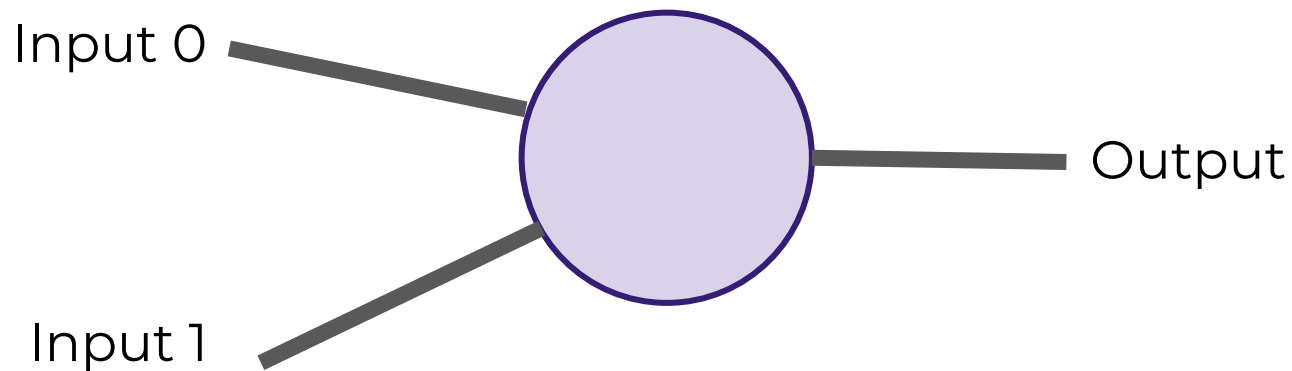
- The artificial neuron also has inputs and outputs!





# Deep Learning


- This simple model is known as a perceptron.





# Deep Learning


- Simple example of how it can work.





# Deep Learning

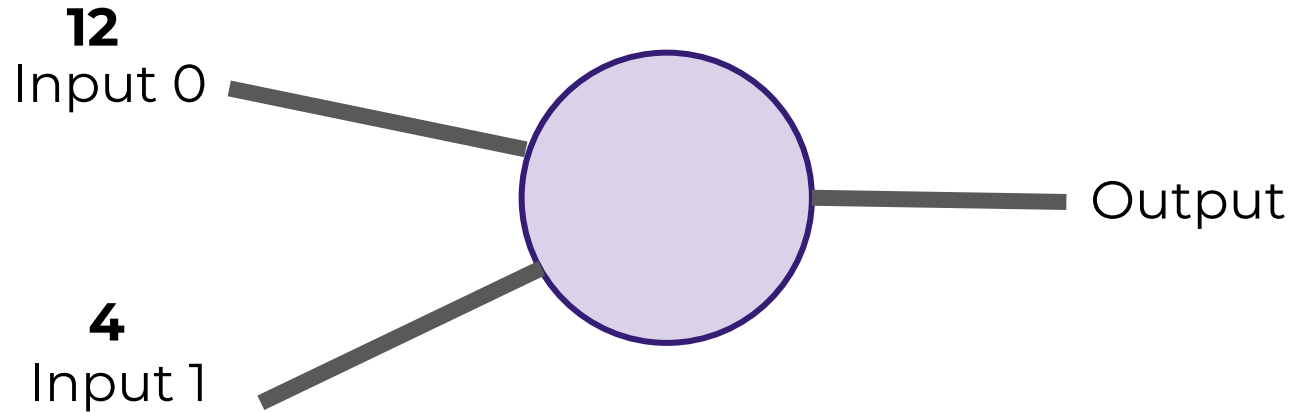
- We have two inputs and an output





# Deep Learning


- Inputs will be values of features





# Deep Learning


- Inputs are multiplied by a weight





# Deep Learning

- Weights initially start off as random








# Deep Learning


- Weights initially start off as random





# Deep Learning


- Inputs are now multiplied by weights





# Deep Learning


- Inputs are now multiplied by weights





# Deep Learning


- Then these results are passed to an activation function.





# Deep Learning


- Many activation functions to choose from, we'll cover this in more detail later!





# Deep Learning


- For now our activation function will be very simple...





# Deep Learning


- If sum of inputs is positive return 1, if sum is negative output 0.





# Deep Learning

- In this case  $6 - 4 = 2$  so the activation function returns 1.








# Deep Learning


- There is a possible issue. What if the original inputs started off as zero?





# Deep Learning


- Then any weight multiplied by the input would still result in zero!





# Deep Learning


- We fix this by adding in a bias term, in this case we choose 1.





# Deep Learning

- So what does this look like mathematically?





# Deep Learning

- Let's quickly think about how we can represent this perceptron model mathematically:

$$\sum_{i=0}^n w_i x_i + b$$



# Deep Learning

- Once we have many perceptrons in a network we'll see how we can easily extend this to a matrix form!

$$\sum_{i=0}^n w_i x_i + b$$



# Deep Learning

- Review
  - Biological Neuron
  - Perceptron Model
  - Mathematical Representation



# Introduction to Neural Networks





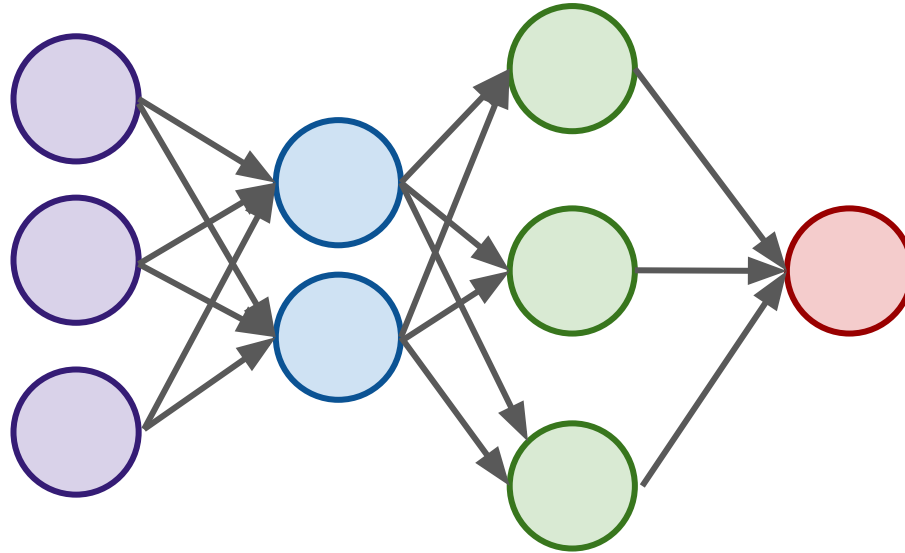
# Deep Learning

- We've seen how a single perceptron behaves, now let's expand this concept to the idea of a neural network!
- Let's see how to connect many perceptrons together and then how to represent this mathematically!



# Deep Learning

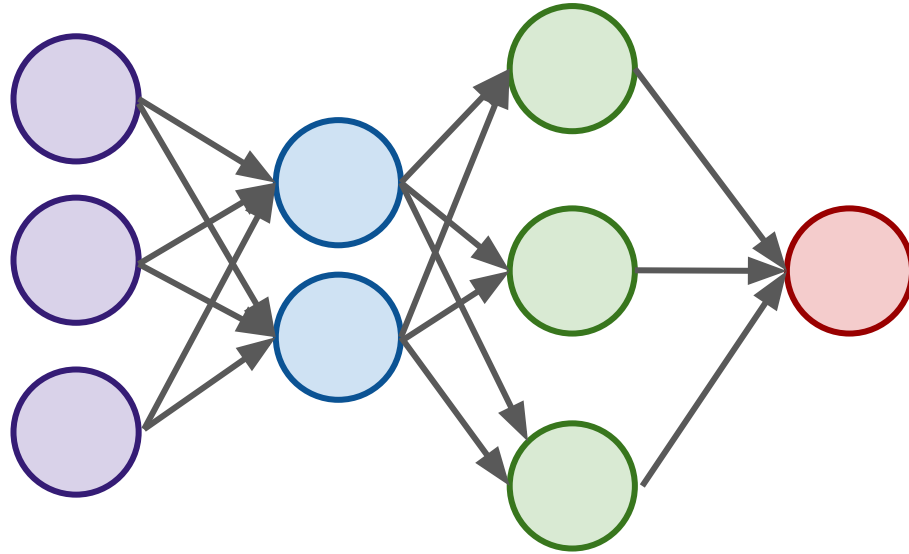
- Multiple Perceptrons Network





# Deep Learning

- Input Layer. 2 hidden layers. Output Layer





# Deep Learning

- Input Layers
  - Real values from the data
- Hidden Layers
  - Layers in between input and output
  - 3 or more layers is “deep network”
- Output Layer
  - Final estimate of the output



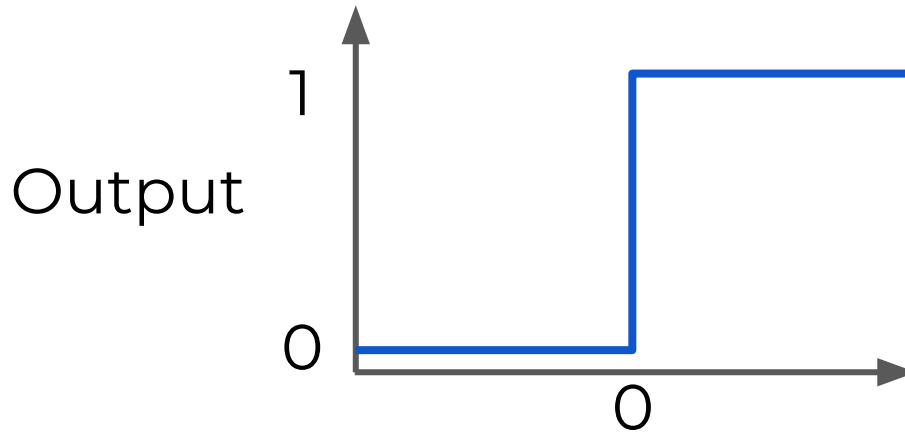
# Deep Learning

- As you go forwards through more layers, the level of abstraction increases.
- Let's now discuss the activation function in a little more detail!



# Deep Learning

- Previously our activation function was just a simple function that output 0 or 1.

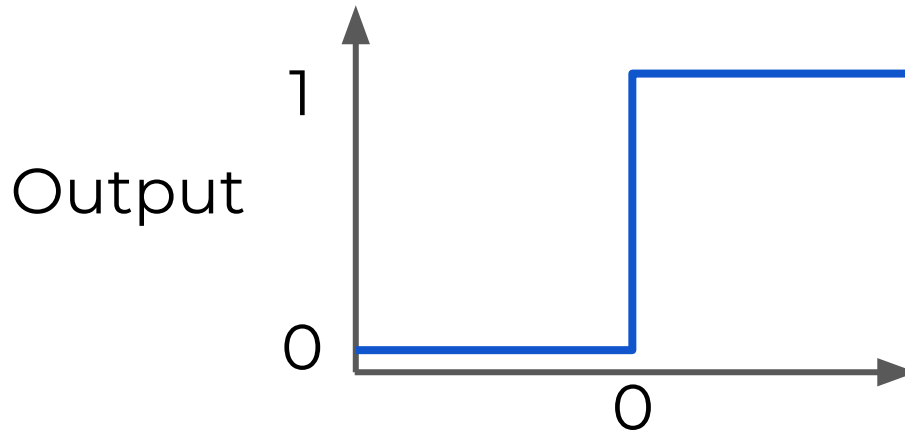


$$z = wx + b$$



# Deep Learning

- This is a pretty dramatic function, since small changes aren't reflected.

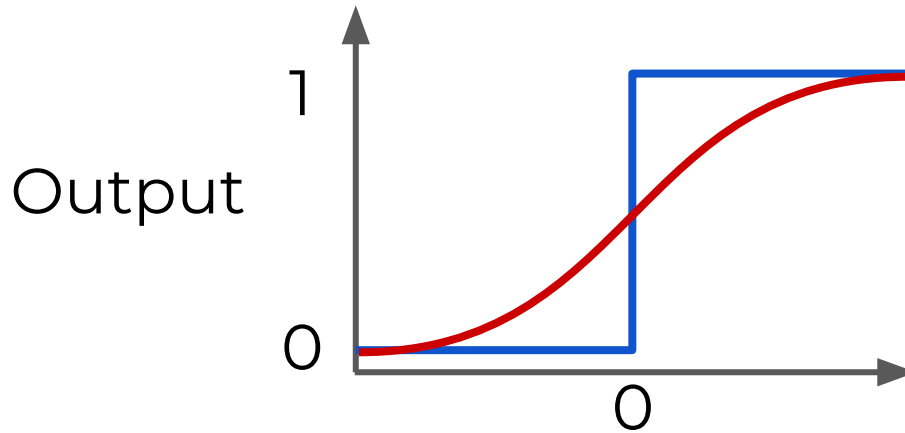


$$z = wx + b$$



# Deep Learning

- It would be nice if we could have a more dynamic function, for example the red line!

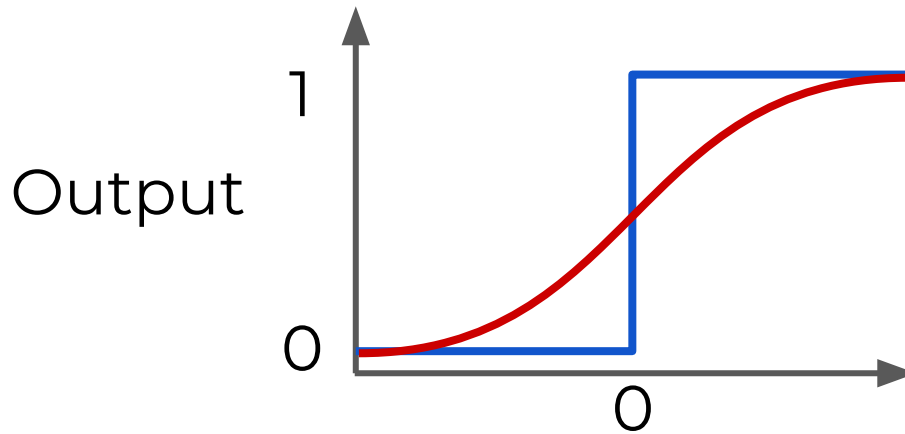






# Deep Learning

- Lucky for us, this is the sigmoid function!




$$f(x) = \frac{1}{1 + e^{-(x)}}$$



# Deep Learning

- Changing the activation function used can be beneficial depending on the task!




$$f(x) = \frac{1}{1 + e^{-(x)}}$$



# Deep Learning

- Let's discuss a few more activation functions that we'll encounter!




$$z = wx + b$$



# Deep Learning

- Hyperbolic Tangent:  $\tanh(z)$



$$z = wx + b$$

$$\cosh x = \frac{e^x + e^{-x}}{2}$$

$$\sinh x = \frac{e^x - e^{-x}}{2}$$


$$\tanh x = \frac{\sinh x}{\cosh x}$$



# Deep Learning

- Rectified Linear Unit (ReLU): This is actually a relatively simple function:  $\max(0, z)$

Output



$$z = wx + b$$



# Deep Learning

- ReLu and tanh tend to have the best performance, so we will focus on these two.
- Deep Learning libraries have these built in for us, so we don't need to worry about having to implement them manually!



# Deep Learning

- As we continue on, we'll also talk about some more state of the art activation functions.
- Up next, we'll discuss cost functions, which will allow us to measure how well these neurons are performing!



# Cost Functions





# Deep Learning

- Let's now explore how we can evaluate performance of a neuron!
- We can use a cost function to measure how far off we are from the expected value.



# Deep Learning

- We'll use the following variables:
  - $y$  to represent the true value
  - $a$  to represent neuron's prediction
- In terms of weights and bias:
  - $w * x + b = z$
  - Pass  $z$  into activation function  $\sigma(z) = a$



# Deep Learning

- Quadratic Cost
  - $C = \Sigma(y-a)^2 / n$
- We can see that larger errors are more prominent due to the squaring.
- Unfortunately this calculation can cause a slowdown in our learning speed.



# Deep Learning

- Cross Entropy
  - $C = (-1/n) \sum (y \cdot \ln(a) + (1-y) \cdot \ln(1-a))$
- This cost function allows for faster learning.
- The larger the difference, the faster the neuron can learn.



# Deep Learning

- We now have 2 key aspects of learning with neural networks, the neurons with their activation function and the cost function.
- We're still missing a key step, actually "learning"!



# Deep Learning

- We need to figure out how we can use our neurons and the measurement of error (our cost function) and then attempt to correct our prediction, in other words, “learn”!



# Deep Learning

- In the next lecture we'll briefly cover how we can do this with Gradient Descent!



# Gradient Descent and Backpropagation





# Deep Learning

- If you've dabbled in machine learning before, you may have already heard of Gradient Descent!
- Let's quickly go over it with a high level overview!




# Deep Learning

- Gradient descent is an optimization algorithm for finding the minimum of a function.
- To find a local minimum, we take steps proportional to the negative of the gradient.



# Deep Learning

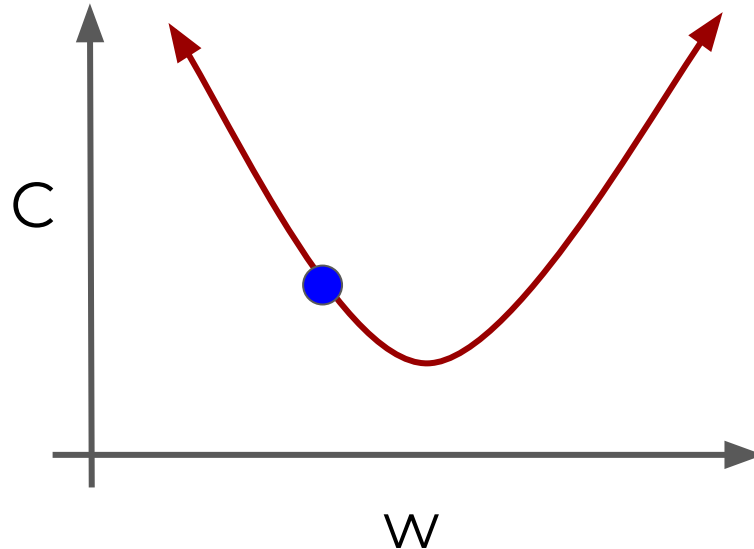
- Gradient Descent (in 1 dimension)





# Deep Learning

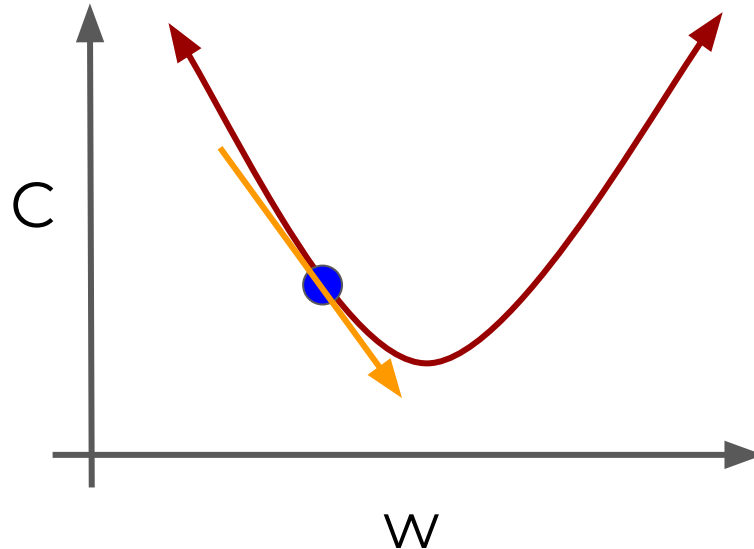
- Gradient Descent (in 1 dimension)





# Deep Learning


- Gradient Descent (in 1 dimension)





# Deep Learning

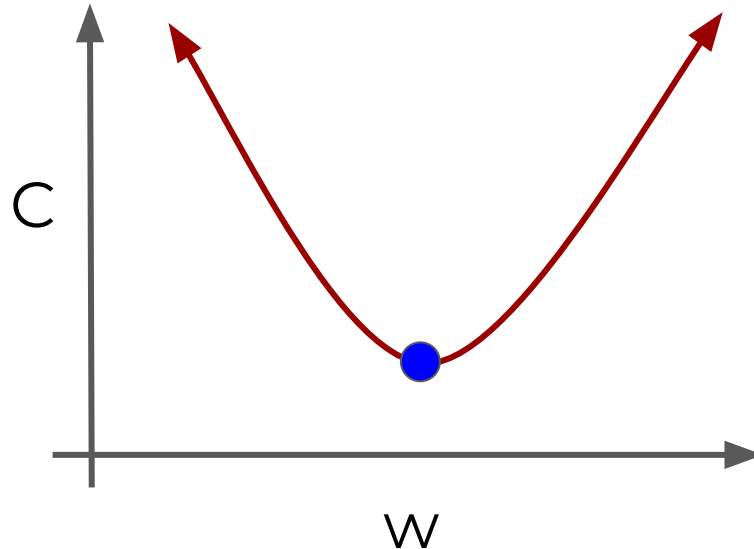
- Gradient Descent (in 1 dimension)





# Deep Learning

- Visually we can see what parameter value to choose to minimize our Cost!





# Deep Learning

- Finding this minimum is simple for 1 dimension, but our cases will have many more parameters, meaning we'll need to use the built-in linear algebra that our Deep Learning library will provide!





# Deep Learning

- Using gradient descent we can figure out the best parameters for minimizing our cost, for example, finding the best values for the weights of the neuron inputs.



# Deep Learning

- We now just have one issue to solve, how can we quickly adjust the optimal parameters or weights across our entire network?
- This is where backpropagation comes in!



# Deep Learning

- Backpropagation is used to calculate the error contribution of each neuron after a batch of data is processed.
- It relies heavily on the chain rule to go back through the network and calculate these errors.



# Deep Learning

- Backpropagation works by calculating the error at the output and then distributes back through the network layers.
- It requires a known desired output for each input value (supervised learning).



# Deep Learning

- The implementation of backpropagation will be further clarified when we dive into the math example!
- For now let's finish off our high level discussion with TensorFlow's playground!



# TensorFlow Playground



# Deep Learning

- Go to:
  - [playground.tensorflow.org](https://playground.tensorflow.org)



# **Manual Neural Network**

## **Part 2 - Operation**





# Deep Learning


- Operation Class
  - Input Nodes
  - Output Nodes
  - Global Default Graph Variable
  - Compute
    - Overwritten by extended classes



# Deep Learning

- Graph - A global variable

Constant



1

Constant



2

Operation






# Deep Learning

- Graph

Constant



1

Constant



2

Add(Operation)




3



# Deep Learning

- Graph

Constant



1

Constant



2

Multiply(Operation)



2



# **Manual Neural Network Variables, Placeholders, and Graphs**



# Deep Learning

- Placeholder - An “empty” node that needs a value to be provided to compute output.
- Variables - Changeable parameter of Graph
- Graph - Global Variable connecting variables and placeholders to operations.



# Let's get started!



# Manual Neural Network Session





# Deep Learning

- Now that the Graph has all the nodes, we need to execute all the operations within a Session.
- We'll use a PostOrder Tree Traversal to make sure we execute the nodes in the correct order.



# Manual Neural Network Classification



# Deep Learning

- $y = mx + b$
- $y = -1x + 5$
- Remember that both  $y$  and  $x$  are features!
- $\text{Feat2} = -1 * \text{Feat1} + 5$
- $\text{Feat2} + \text{Feat1} - 5 = 0$
- $\text{FeatMatrix}[1, 1] - 5 = 0$